



# Real World Kernel Pool Exploitation

Kostya Kortchinsky

Immunity, Inc.

*SyScan'08 Hong Kong*



- Introduction
- The Kernel Pool and its structures
- Allocation and Free algorithm
- Exploiting a Kernel Pool overflow
- MS08-001: IGMPv3 Kernel Pool overflow
- Conclusion



KNOWING YOU'RE SECURE

# Introduction



# Introduction

KNOWING YOU'RE SECURE

- Kernel vulnerabilities are more popular
  - Less protections than in userland
    - Kernel allocations are highly performance optimized
    - Less room for heap cookies and other roadblocks
  - Code has had less attention than userland services
- Several kernel pool overflows to exploit
  - Our example: MS08-001 (IGMPv3 Overflow)

# Population

KNOWING YOU'RE SECURE

People who write exploits

People who write  
Windows  
overflows

People who  
write Windows  
Kernel Pool  
Overflows



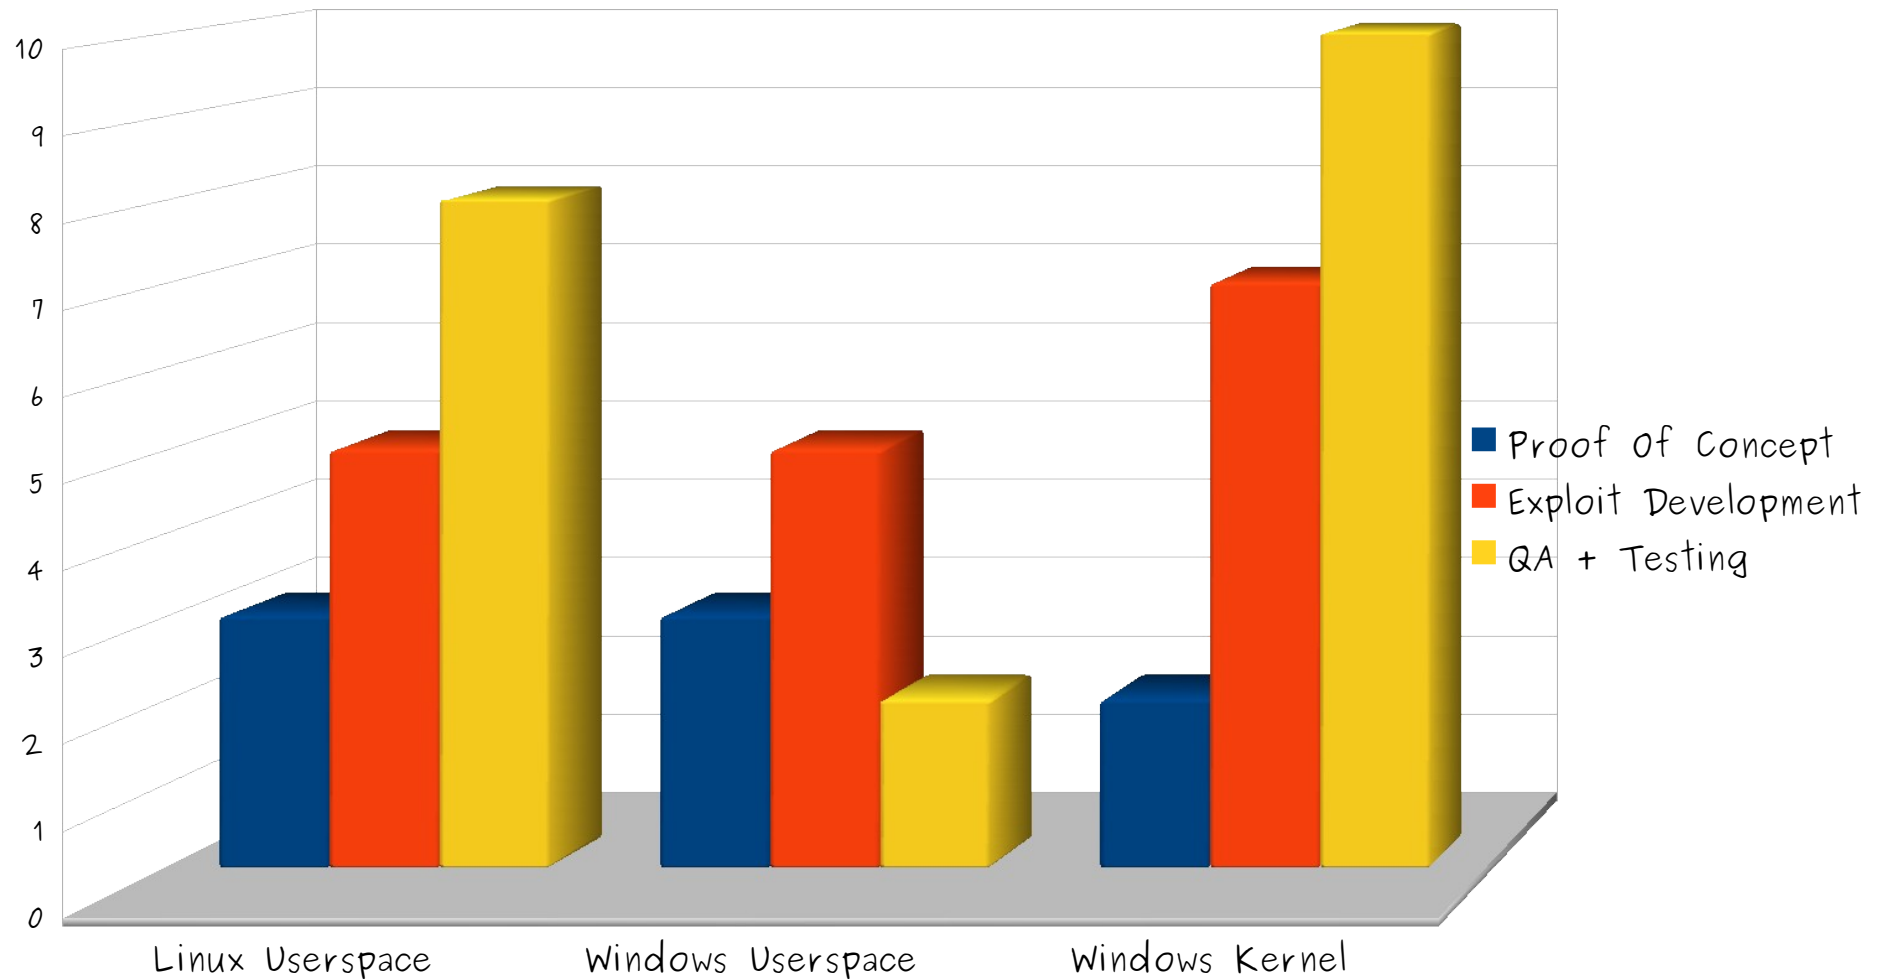


# Other considerations

KNOWING YOU'RE SECURE

- Because few people know how to write them, Windows Kernel Heap overflows are often mischaracterized as “Denial of Service”
- Huge investment in exploit development before it is known if a reliable exploit can be created for any bug
  - If it costs 100K dollars to even know if you have a risk from any particular vulnerability, can you afford to find out?

# Diversity increases QA costs dramatically





# Addresses May Vary

KNOWING YOU'RE SECURE

- The following slides assume:
  - Structures, constants and offsets are from ntkrnlpa.exe 5.1.2600.2180
  - nt!KeNumberNodes is 1 (non NUMA architecture)
- *Reminder:* Windows XP SP2 has 4 kernels
  - ntoskrnl.exe
  - ntkrnlpa.exe: PAE
  - ntkrnlmp.exe: MP
  - ntkrnpamp.exe: PAE+MP (NUMA aware)





# The Kernel Pool and its structures



# Kernel Pool vs. Userland Heap

KNOWING YOU'RE SECURE

- Quite similar
- Only a few pools for all the kernel allocations
  - Think LsaSs default heap
- Kernel pool is designed to be fast
  - As few checks as possible
    - No kernel pool cookie
    - No kernel pool safe unlink



# Kernel Pool

KNOWING YOU'RE SECURE

- Used by Windows for dynamic memory allocations within kernel land by functions:
  - nt!ExAllocatePool, nt!ExAllocatePoolWithTag, ...
  - nt!ExFreePool, nt!ExFreePoolWithTag, ...
- There are several kernel pools, default being:
  - One *non-paged* pool
  - Two *paged* pools
  - One *session paged* pool
- Pools are defined thanks to structure nt!\_POOL\_DESCRIPTOR, and stored in nt!PoolVector



# Non Paged Pool

KNOWING YOU'RE SECURE

- *Non pageable* system memory
- Can be accessed from any IRQL
- Scarce resource
- Descriptor for non paged pool is static:
  - nt!NonPagedPoolDescriptor in .data section
- It is initialized in nt!InitializePool
- Pointer to the descriptor is stored in entry 0 of nt!PoolVector





- *Pageable* system memory
- Number of paged pools defined by nt!  
ExpNumberOfPagedPools (default is **2**)
- An array of pool descriptors is dynamically allocated and initialized in nt!InitializePool
  - NonPagedPool allocation
  - One more descriptor than number of paged pools
- Pointer to array stored in entry 1 of nt!  
PoolVector and in nt!ExpPagedPoolDescriptor



# Session Paged Pool

KNOWING YOU'RE SECURE

- *Pageable* system memory
- Descriptor for the session paged pool is located in the session space
  - PagedPool member of nt!MmSessionSpace structure
    - Usually  $0\text{xbf}7\text{f}0000 + 0\text{x}244$
- Initialized in nt!MiInitializeSessionPool
  - used for session space allocations
  - do not use lookaside lists
- Pointer to the descriptor stored in nt!ExpSessionPoolDescriptor

# nt!\_POOL\_DESCRIPTOR

KNOWING YOU'RE SECURE

```
kd> dt nt!_POOL_DESCRIPTOR
+0x000 PoolType           : _POOL_TYPE
+0x004 PoolIndex          : Uint4B
+0x008 RunningAllocs      : Uint4B
+0x00c RunningDeAllocs    : Uint4B
+0x010 TotalPages         : Uint4B
+0x014 TotalBigPages      : Uint4B
+0x018 Threshold          : Uint4B
+0x01c LockAddress        : Ptr32 Void
+0x020 PendingFrees       : Ptr32 Void
+0x024 PendingFreeDepth   : Int4B
+0x028 ListHeads          : [512] _LIST_ENTRY
```

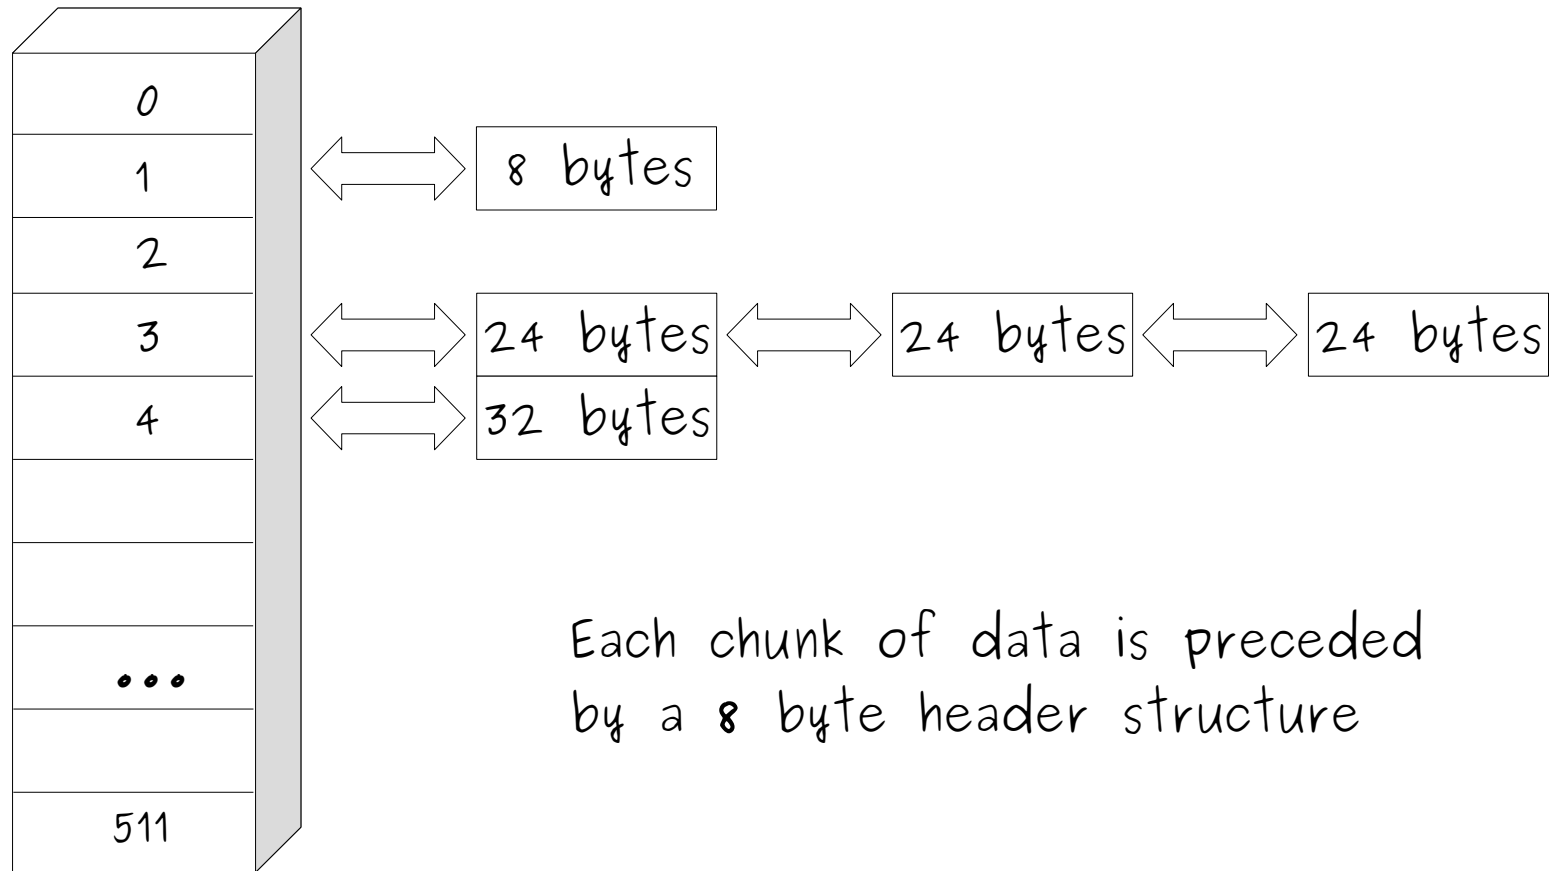


- Some of the members:
  - *PoolType*: NonPagedPool=0, PagedPool=1, ...
  - *PoolIndex*: 0 for non-paged pool and paged session pool, index of the pool descriptor in nt!ExpPagedPoolDescriptor for paged pools
  - *ListHeads*: 512 double-linked lists of free memory chunks of the same size (8 byte granularity)
    - List number for a requested allocation size is calculated by  $\text{BlockSize} = (\text{NumberOfBytes} + 0xf) \gg 3$
    - Thus can be used for allocations up to 0xff0 (4080) bytes



# ListHeads

KNOWING YOU'RE SECURE



# nt!\_POOL\_HEADER

KNOWING YOU'RE SECURE

```
kd> dt nt!_POOL_HEADER
+0x000 PreviousSize      : Pos 0, 9 Bits
+0x000 PoolIndex         : Pos 9, 7 Bits
+0x002 BlockSize         : Pos 0, 9 Bits
+0x002 PoolType          : Pos 9, 7 Bits
+0x000 Ulong1            : Uint4B
+0x004 ProcessBilled     : Ptr32 _EPROCESS
+0x004 PoolTag           : Uint4B
+0x004 AllocatorBackTraceIndex : Uint2B
+0x006 PoolTagHash       : Uint2B
```



# Chunk Header

KNOWING YOU'RE SECURE

- Explanations on some of the members:
  - *PreviousSize*: *BlockSize* of the preceding chunk
    - 0 if chunk is located at the beginning of a page
  - *BlockSize*: Number of bytes requested plus header size rounded up to a multiple of 8, divided by 8
    - or  $(\text{NumberOfBytes} + 0xf) \gg 3$
  - *PoolIndex*: same definition as for the descriptor
  - *PoolType*: 0 if free,  $(\text{PoolType} + 1) | 4$  if allocated
  - *PoolTag*: Usually 4 printable characters identifying the code responsible for the allocation



# Free Chunk Header

KNOWING YOU'RE SECURE

- When *PoolType*=0, the chunk header is followed by a `nt!_LIST_ENTRY` structure
  - This is the entry pointed to by the ListHeads double linked list
- Chunks freed to the lookaside lists remain the same, their *PoolType* is non 0

```
kd> dt nt!_LIST_ENTRY
```

+0x000 Flink	: Ptr32 _LIST_ENTRY
+0x004 Blink	: Ptr32 _LIST_ENTRY





- Like for userland heaps, kernel uses lookaside lists for faster allocating and freeing of small chunks of data
  - Maximum *BlockSize* being 32 (or 256 bytes)
- They are defined in the processor control block
  - 32 entry `PPPagedLookasideList`
  - 32 entry `PPNPagedLookasideList`
- Each entry holds **2** single chained lists of nt! `_GENERAL_LOOKASIDE` structures: one “per processor” *P*, one “system wide” *L*

# nt!\_GENERAL\_LOOKASIDE

KNOWING YOU'RE SECURE

```
kd> dt nt!_GENERAL_LOOKASIDE
+0x000 ListHead          : _SLIST_HEADER
+0x008 Depth             : Uint2B
+0x00a MaximumDepth      : Uint2B
+0x00c TotalAllocates    : Uint4B
+0x010 AllocateMisses    : Uint4B
+0x010 AllocateHits      : Uint4B
+0x014 TotalFrees        : Uint4B
+0x018 FreeMisses        : Uint4B
+0x018 FreeHits          : Uint4B
+0x01c Type              : _POOL_TYPE
+0x020 Tag                : Uint4B
+0x024 Size              : Uint4B
+0x028 Allocate          : Ptr32      void*
+0x02c Free               : Ptr32      void
+0x030 ListEntry          : _LIST_ENTRY
+0x038 LastTotalAllocates : Uint4B
+0x03c LastAllocateMisses : Uint4B
+0x03c LastAllocateHits  : Uint4B
+0x040 Future             : [2] Uint4B
```

# nt!MmNonPagedPoolFreeListHead

KNOWING YOU'RE SECURE

- Static array of 4 double linked lists for non-paged free chunks bigger than a page
    - Index in the array is (SizeOfChunkInPages-1)
      - Last entry also has everything bigger than 4 pages
  - Structure used is nt!\_MMFREE\_POOL\_ENTRY
- ```
kd> dt nt!_MMFREE_POOL_ENTRY
+0x000 List           : _LIST_ENTRY
+0x008 Size           : Uint4B
+0x00c Signature      : Uint4B
+0x010 Owner          : Ptr32 _MMFREE_POOL_ENTRY
```
- Thus free non-paged “big” chunks are linked through the pages themselves



# Allocation and Free algorithms





# Allocation Summary

KNOWING YOU'RE SECURE

- Windows kernel pool allocates in small chunks:
  - 8 byte granularity up to 4080 bytes (included)
    - Used to be 32 for Windows 2000
  - Page granularity above
- Makes extensive use of optimized lists:
  - Single linked Lookaside lists up to 256 bytes
  - Double linked ListHeads lists up to 4080 bytes
- Splits an entry if a chunk of the exact size cannot be found
- Expands the pool if needed

# Simplified Allocation Algorithm

## nt!ExAllocatePoolWithTag (1/2)

- If NumberOfBytes > 0xffff:
  - Call nt!MiAllocatePoolPages
- If PagedPool requested:
  - If BlockSize ≤ 0x20:
    - Try the “per processor” paged lookaside list
    - If failed, try the “system wide” paged lookaside list
    - Return on success
  - Try and lock a paged pool descriptor
- Else:
  - If BlockSize ≤ 0x20:
    - Try the “per processor” non-paged lookaside list

# Simplified Allocation Algorithm

## nt!ExAllocatePoolWithTag (2/2)

- If failed, try the “system wide” non-paged lookaside list
- Return on success
- Try and lock the non-paged pool descriptor
- Use ListHeads of currently locked pool:
  - Use 1<sup>st</sup> non empty ListHeads[n]
    - With  $\text{BlockSize} \leq n < 512$
    - Split entry if bigger than needed
    - Return on success
  - If failed, expand the pool by adding a page
    - Try again!



# Free Chunk Splitting

KNOWING YOU'RE SECURE


- If the algorithm has selected a chunk larger than the requested `NumberOfBytes`, it is split
  - If the chunk is at the start of a page:
    - Take the allocation from the *front* of the chunk
  - Otherwise:
    - Take the allocation from the *end* of the chunk
- The remaining part is inserted in the correct list




# Splitting Schema

KNOWING YOU'RE SECURE

Free chunk at the beginning of a page,  
allocated chunk goes at the front



Otherwise, allocated chunk goes at the end





- Free algorithm works pretty much as expected:
  - It will use Lookaside lists for chunks up to 256 bytes
    - If they are not full already
  - It will use ListHeads for chunks up to 4080 bytes
- Merges contiguous free chunks to lower fragmentation
- Releases pages if necessary

# Simplified Free Algorithm

nt!ExFreePoolWithTag (1/2)

- If P is page aligned:
  - Call nt!MiFreePoolPages
- If BlockSize ≤ 0x20:
  - If PoolType = PagedPool:
    - Put in “per processor” paged lookaside list
    - If failed, put in “system wide” paged lookaside list
    - Return on success
  - Else:
    - Put in “per processor” non-paged lookaside list
    - If failed, put in “system wide” non-paged lookaside list
    - Return on success

# Simplified Free Algorithm



nt!ExFreePoolWithTag (2/2)

- If next chunk is free and not page aligned:
  - Merge with current chunk
- If previous chunk is free:
  - Merge with current chunk
- If resulting chunk is a full page:
  - Call nt!MiFreePoolPages
- Else:
  - Add chunk to the tail of the correct ListHeads
    - Based on PoolType, PoolIndex and BlockSize of chunk






Chunk being freed



Merge #1



Merge #2





# Exploiting a Kernel Pool overflow



- Discrepancies in the kernel pool will most likely result in a BugCheck (Blue Screen)

0x19: BAD\_POOL\_HEADER

0x41: MUST\_SUCCEED\_POOL\_EMPTY

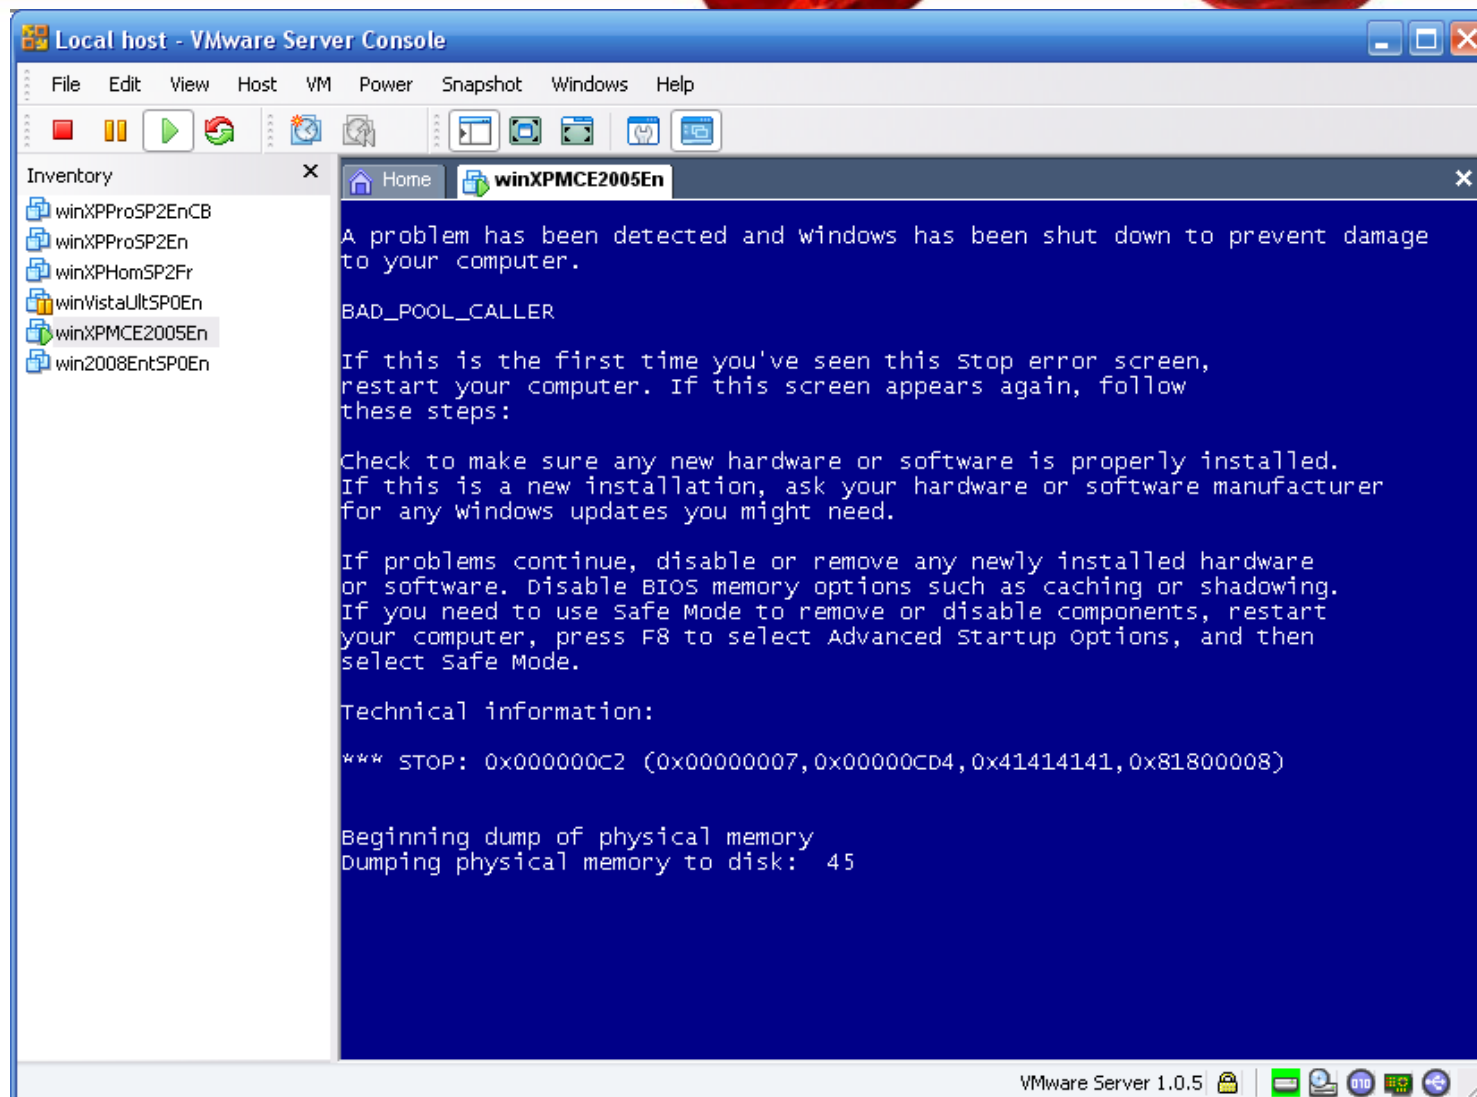
0xc1: SPECIAL\_POOL\_DETECTED\_MEMORY\_CORRUPTION

0xc2: BAD\_POOL\_CALLER

- Some are only present in Checked Build
- Avoid those when exploiting a pool overflow!

# BugCheck Example

KNOWING YOU'RE SECURE





# Some BugCheck Conditions

KNOWING YOU'RE SECURE

- nt!ExFreePoolWithTag:
  - BugCheck 0xc2, 7 if *PoolType*&4=0
    - The chunk attempted to be freed is already free
  - BugCheck 0x19, 0x20 if *PreviousSize* of next chunk is  $\neq$  *BlockSize* of current chunk
- Checked Build:
  - BugCheck 0x19, 3 if (Entry→Flink)→Blink! $\neq$ Entry or (Entry→Blink)→Flink $\neq$ Entry
    - It didn't make it to retail build, thanks Microsoft!



Yes!

- Unlike in userland heaps, there is no such thing as a kernel pool *cookie*
- There is no safe unlinking in retail build



# Kernel Pool Unlink

KNOWING YOU'RE SECURE

- Removing an entry 'e' from a double linked list:

```
PLIST_ENTRY b,f;  
f=e→Flink;  
b=e→Blink;  
b→Flink=f;  
f→Blink=b;
```

- This leads to a usual write4 primitive:

```
*(where)=what  
*(what+4)=where
```

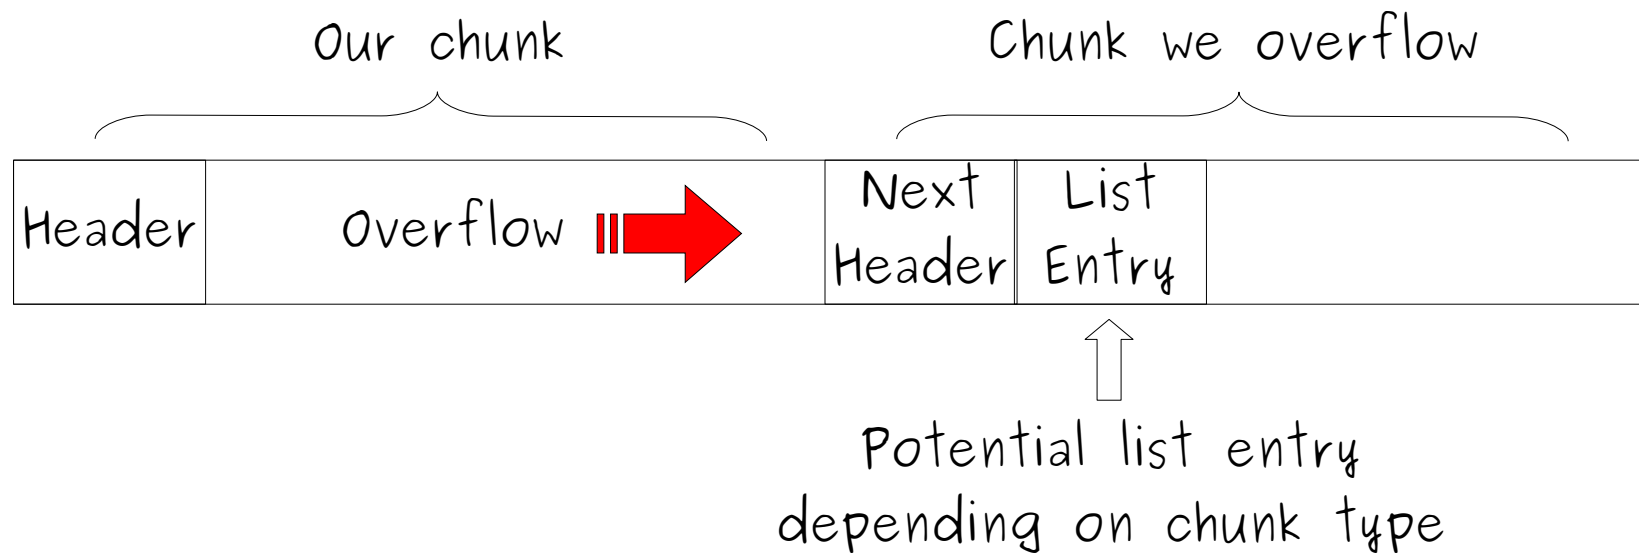
- Locating the unlinks:

- nt!ExFreePoolWithTag: when merging chunks
- nt!MiAllocatePoolPages, nt!MiFreePoolPages, ...

# Notations

KNOWING YOU'RE SECURE

## Kernel Pool Overflow







- Summary list of write4 techniques:
  - Write4 on Merge with Next
    - When freeing our chunk
  - Write4 on Merge with Previous
    - When freeing the chunk we overflowed
  - Write4 on *ListHeads* Unlink
    - If we overflowed an entry in a *ListHeads* list
  - Write4 on MmNonPagedPoolFreeListHead Unlink
    - If we overflowed an entry in a MmNonPagedPoolFreeListHead list

# Write4 on Merge with Next

KNOWING YOU'RE SECURE

Case #1

- When our chunk is freed:
  - If *PreviousSize* of next chunk is = *BlockSize* of current chunk
    - To avoid BugCheck 0x19
  - If *BlockSize* > 0x20 (or lookaside lists are full)
    - To avoid a free to lookaside
  - If *PoolType* of next chunk is 0
  - If *BlockSize* of next chunk is > 1
    - Otherwise it means there is no list entry
  - Then merge with next chunk:
    - And **unlink** happens on list entry of next chunk



- Allocate a chunk of size 256-4080 bytes
  - Works with smaller chunks if lookaside lists are full
- Craft a free header after our chunk with:
  - Correct *PreviousSize*
  - *PoolType*=0
  - Wanted *Flink* and *Blink*
  - Requires a minimum overflow of about 16 bytes
- Write4 happens when our allocated chunk is freed

# Write4 on Merge with Previous

KNOWING YOU'RE SECURE

Case #2

- When the chunk we overflowed is freed:
  - If *PreviousSize* of next chunk is = *BlockSize* of current chunk
    - To avoid BugCheck 0x19
  - If *BlockSize* > 0x20 (or lookaside lists are full)
    - To avoid a free to lookaside
  - If *PoolType* of previous chunk is 0
  - If *BlockSize* of previous chunk is > 1
    - Otherwise it means there is no list entry
  - Then merge with previous chunk:
    - And **unlink** happens on list entry of previous chunk





# Exploit Case #2 (1/2)


KNOWING YOU'RE SECURE

- Allocate a chunk
- Next chunk **must** be allocated and of size 256-4080 bytes (not as easy as it looks)
  - Works with smaller chunks if lookaside lists are full
- Craft a fake free header and list entry at the end of our chunk (with realistic sizes)
- Overflow *PreviousSize* of next chunk to make it point to our fake header
  - Requires a ~1 byte overflow!
    - 0x00 can work if *PreviousSize* > 0x100


# Exploit Case #2 (2/2)

KNOWING YOU'RE SECURE

- A bigger overflow would require knowledge of *BlockSize* of next chunk to avoid BugCheck
  - Or enough bytes to craft another header after (~257)
- Write4 happens when next chunk is freed
  - Also works with page aligned chunks!



Overflow PreviousSize of next chunk





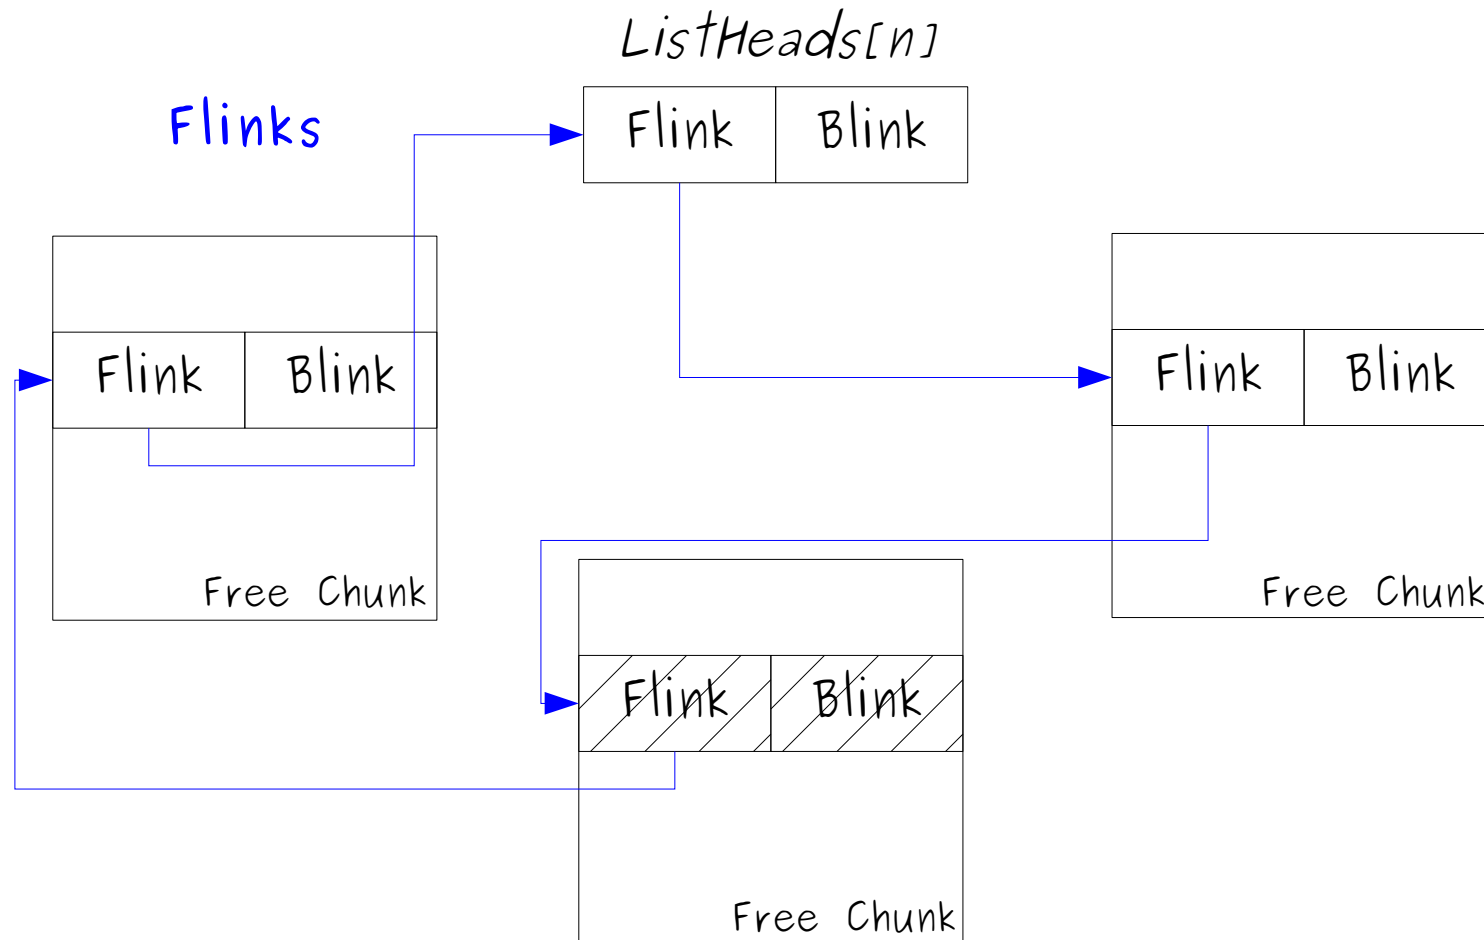
# ListHeads Write4

KNOWING YOU'RE SECURE

- When the chunk we overflowed is allocated:
  - If the chunk was requested through ListHeads list
    - No other constraint on *BlockSize*, *PreviousSize*, ...
  - Then the overflowed list entry is removed from the ListHeads list:
    - We overwrite of ListHeads[BlockSize]→Flink with a pointer we control
- Next time a chunk of *BlockSize* is requested, our pointer will be returned
- **Variations** of the write4 exist based on operations done on the neighboring entries

# ListHeads Illustrated (1/3)

KNOWING YOU'RE SECURE

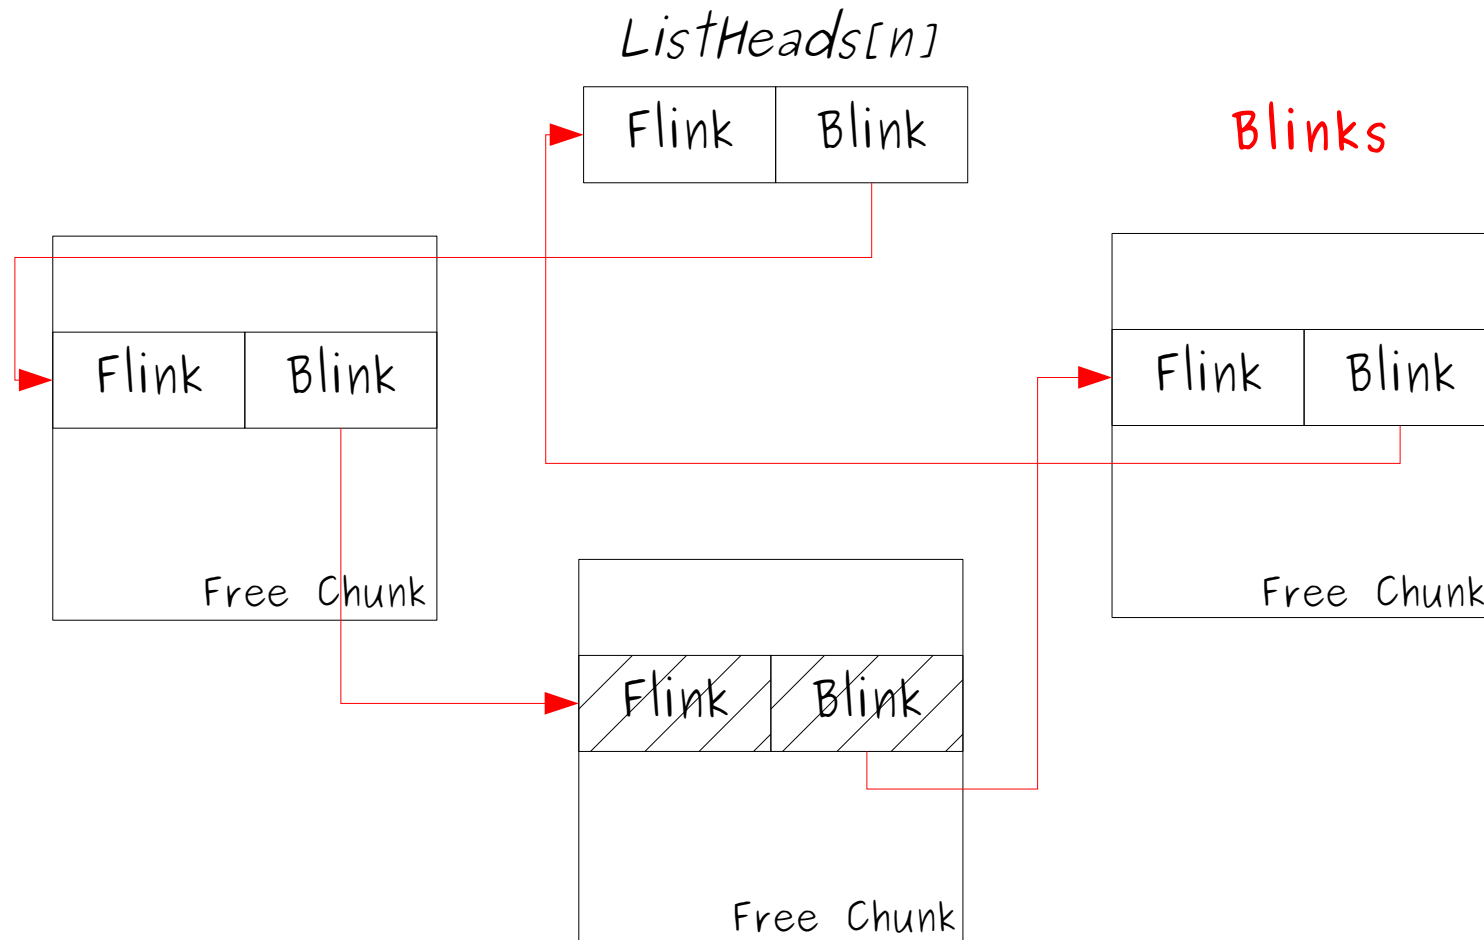


 : overflowed list entry



# ListHeads Illustrated (1/3)

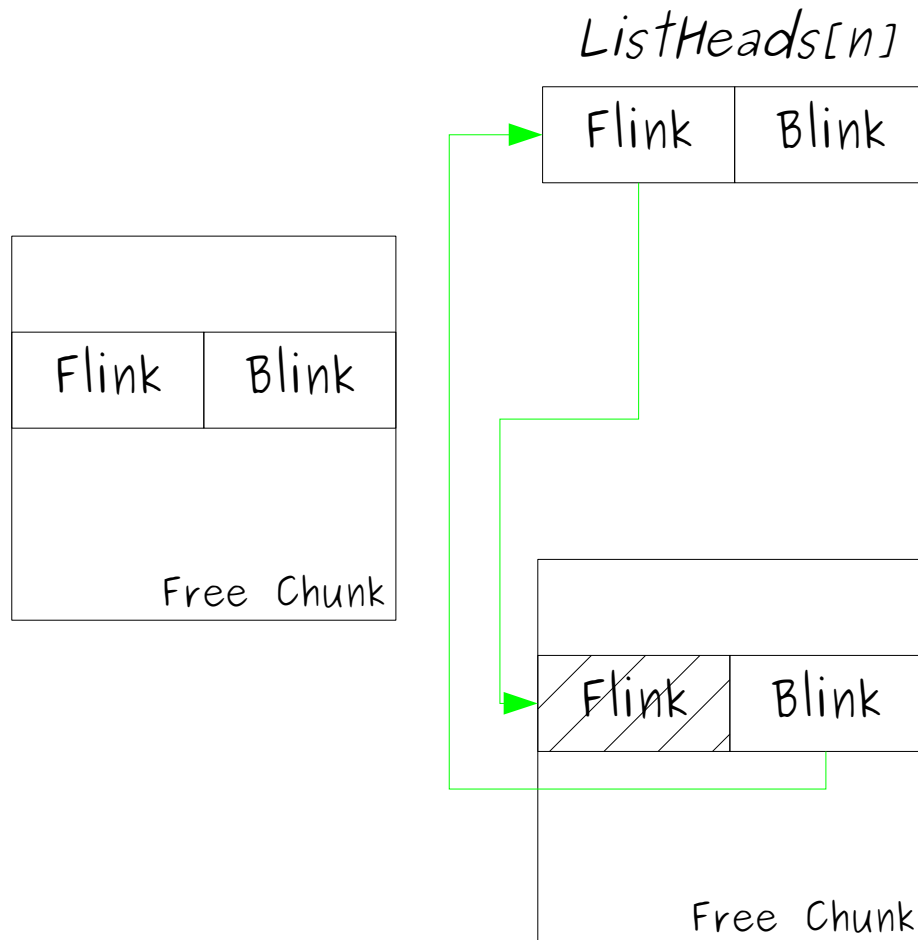
KNOWING YOU'RE SECURE




 : overflowed list entry

# ListHeads Illustrated (2/3)

KNOWING YOU'RE SECURE



Allocation of size  $n$   
unlinks  $\text{ListHeads}[n] \rightarrow \text{Flink}$





```
PLIST_ENTRY b, f;  
f = ListHeads[n] → Flink → Flink;  
b = ListHeads[n] → Flink → Blink;  
b → Flink = f;  
f → Blink = b;
```


# ListHeads Illustrated (3/3)

KNOWING YOU'RE SECURE

*ListHeads[n]*



Allocation of size n  
unlinks  $\text{ListHeads}[n] \rightarrow \text{Flink}$



```
PLIST_ENTRY b,f;  
f=ListHeads[n]→Flink→Flink;  
b=ListHeads[n]→Flink→Blink;  
b→Flink=f;  
f→Blink=b;    ↩ might AV
```

$\text{ListHeads}[n] \rightarrow \text{Flink}$  is now under our control!

# MMFREE\_POOL\_ENTRY Write4

KNOWING YOU'RE SECURE

- When the chunk we overflowed is allocated:
  - If the chunk was requested through nt!  
MmNonPagedPoolFreeListHead list
    - Overflowed chunk is obviously page aligned
  - Then the write4 happens when the  
MMFREE\_POOL\_ENTRY structure is removed from  
the double linked list
- Variants of the write4 exist based on operations  
done on the neighboring list entries





- Most complicated with Kernel Pool overflows:
  - Only a few pools for **all** kernel allocations
    - Lot of allocation and free
  - 4 different kernels based on architecture:
    - Single processor: *ntoskrnl.exe*, *ntkrnlpa.exe*
    - Multi processors: *ntkrnlmp.exe*, *ntkrpamp.exe*
  - A lot of kernel patches for every service pack
    - ⇒ Addresses tend to change, a lot
- Any access violation will most likely end up in a BugCheck ☹



# Some Ideas

KNOWING YOU'RE SECURE

Non Exhaustive

- nt!KiDebugRoutine function pointer
  - Called by nt!KiDispatchException if not NULL
  - Should be kept as last resort
- Context specific function pointers
  - tcpip!tcpxsum\_routine is called right after the overflow in the case of MS08-001
- Function pointers arrays
  - nt!HalDispatchTable
  - Interrupt Dispatch Table (IDT)
- Kernel instructions – page is RWE!

# Write4 into the Kernel

KNOWING YOU'RE SECURE

- Before

```
mov eax, [edx+8]
mov ebx, [edx+0Ch]
mov [ebx], eax
mov [eax+4], ebx
```

Edx points to something we control

```
loc_80543F0B: ; CODE XREF: ExFreePoolWithTag(x,x)+518j
movzx edx, word ptr [edx+2]
```

- After

```
mov eax, [edx+8]
mov ebx, [edx+0Ch]
mov [ebx], eax
mov [eax+4], ebx
```

Jmp edx being 2 bytes long, we can pick the upper 2 so that the write4 doesn't trigger an access violation

```
loc_80543F0B: ; CODE XREF: ExFreePoolWithTag(x,x)+518j
jmp edx
```



# Fixing the Kernel Pool

KNOWING YOU'RE SECURE

- Check for inconsistencies and fix them:
  - Lookaside lists, both “per processor” and systemwide
    - Zero out *Sequence*, *Depth* and *Next* of *ListHead* member of the given `nt!_GENERAL_LOOKASIDE` entry
      - In fact the first 8 bytes
  - ListHeads lists for involved pool descriptor(s)
    - Set `ListHeads[BlockSize]→Flink` and `ListHeads[BlockSize]→Blink` to `&ListHeads[BlockSize]`
  - `nt!MmNonPagedPoolFreeListHead` array





# MS08-001: IGMPv3 Kernel Pool overflow

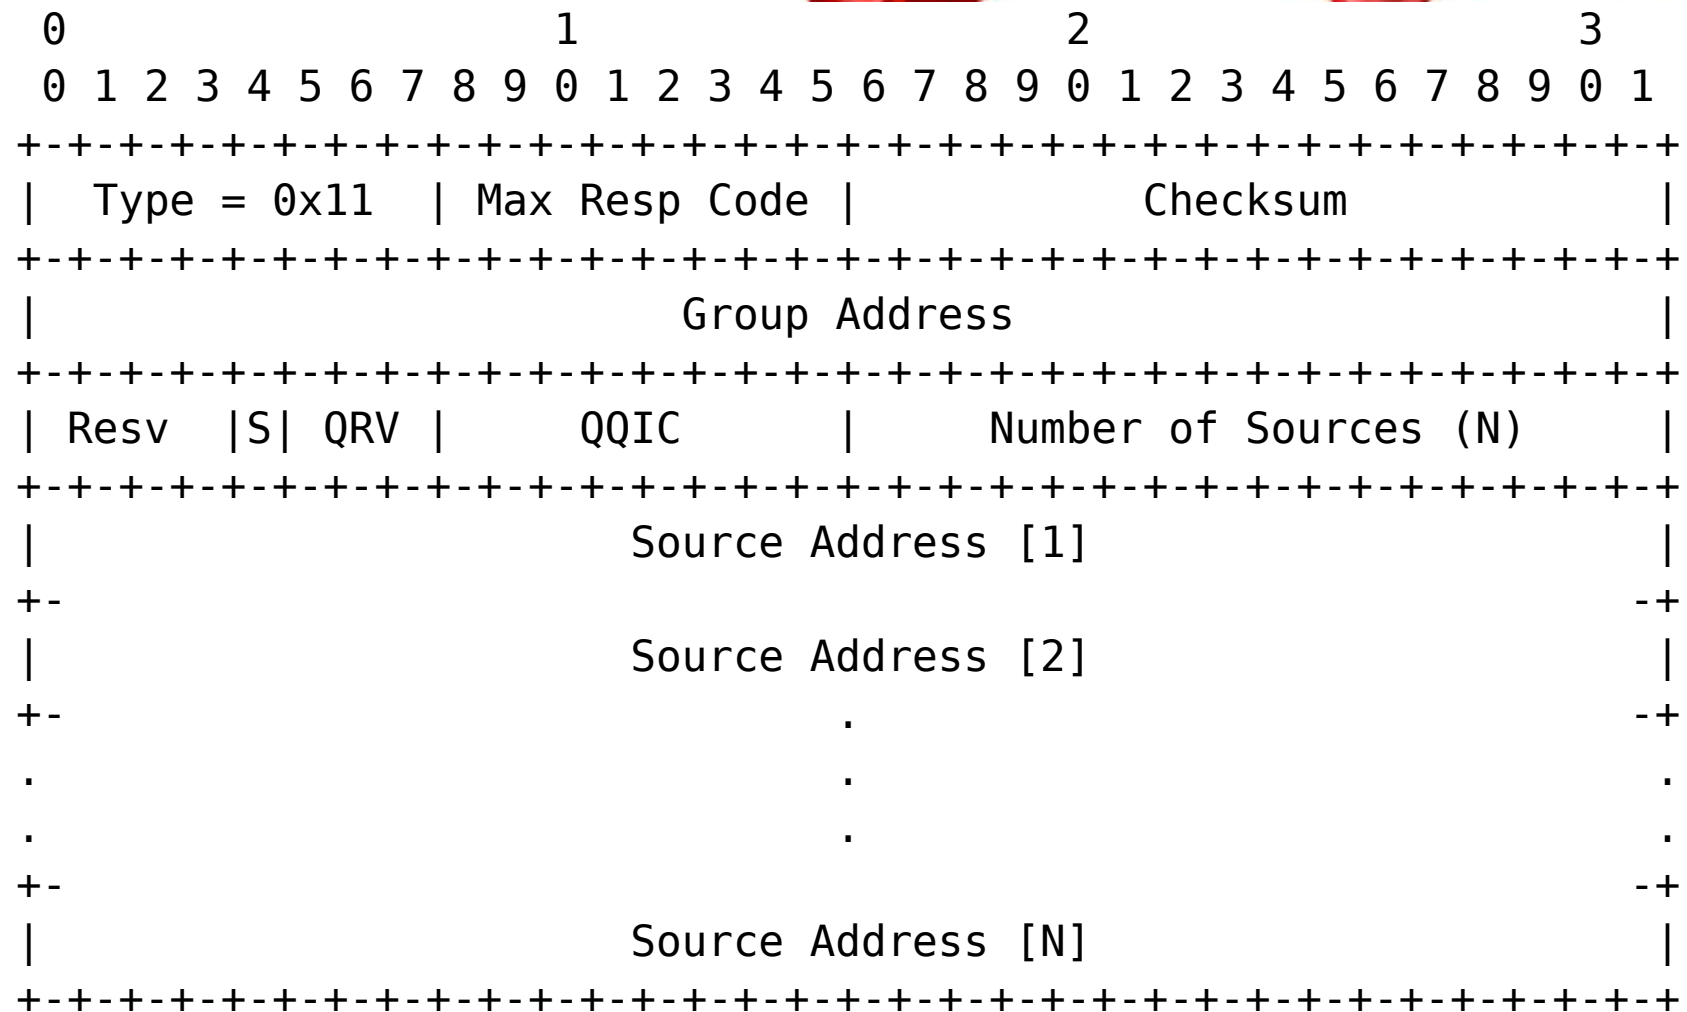


- Remote default kernel pool overflow in Windows XP when parsing IGMPv3 packets
  - Even bypasses default firewall rules!
- Released January, 8<sup>th</sup> 2008 in [MS08-001](#)
  - Along with MLDv2 vulnerability for Vista
- Reported by Alex Wheeler and Ryan Smith of [IBM – Internet Security Systems](#)
- Considered by Microsoft [SWI](#) as “unlikely” exploitable

# IGMPv3 Membership Queries

KNOWING YOU'RE SECURE

RFC 3376





1. Walk a single linked list to count the number of entries (using a 32 bit counter)
2. Allocate some memory:

```
loc_44197:      ; CODE XREF: GetGSIsInRecord(x,x)+18j
push 10h ; Priority
movzx eax, dx
push 'qICT' ; Tag
lea eax, ds:8[eax*4]
push eax ; NumberOfBytes
push ebx ; PoolType
call ds:__imp__ExAllocatePoolWithTagPriority@16
```

❖ Failed!
3. Copy the list entries in the allocated array by walking the list ➡ Overflow





- Send multiple IGMPv3 membership queries
  - Group address has to be valid
  - Can be sent to:
    - Unicast IP address (ie. 10.10.10.42)
    - Broadcast IP address (ie. 10.10.10.255)
    - Multicast IP address (ie. 224.0.0.1)
  - Total number of **unique** source addresses must be 65536 or greater
    - IP addresses in the 224-239 range are ignored
- Wait for the IGMPv3 report to be triggered



- Sending a lot of IGMPv3 membership queries induces *high CPU utilization*
  - Each new source address triggers the allocation of a **20 (0x14) byte** structure
  - The linked list of structures is walked before adding a new element to check for uniqueness of IP ( $O(n^2)$ )
- High CPU usage leads to potential packets dropping
  - Final buffer will not be as expected

## Issues (2/2)

KNOWING YOU'RE SECURE

- IGMPv3 reports are on a *random* timer
  - Can be triggered before all the queries are sent
  - Buffer will not be filled as intended
- 16 bit wrap *usually* means **huge** overflow
  - 65536 entries are put in a 0 byte buffer



- Exponential delay between packets sending
  - Still fast enough to avoid a random report trigger
- Report trigger can be forced if  $QQIC=0$ 
  - Empty the list before the attack
  - Trigger the overflow when all the packets are sent
- Index for the copy is on 16 bit register
  - Overflow size is actually (for  $n > 0xffff$ ):
$$(0x10000 - (n \% 0x10000)) * 4$$
  - For  $n = 0x1fffe$ , allocation size is **0x40000** (page aligned allocation), overflow size is 8






- Pros
  - Relatively small number of packets ( $\sim 180$ )
  - Since the allocated buffer is freed after the overflow, “Write4 on Merge with Next” will always happen
    - Next chunk is easily craftable
  - We control *BlockSize* to some extent, which is interesting for the “Write4 into Kernel” technique
- Cons
  - Overflow is huge! ( $> 0x3f000$  bytes)
    - A large chunk of the kernel pool will be trashed ☹
    - Lot of pool fixing to do

# Why would it work?

KNOWING YOU'RE SECURE


Kernel Pool is filled with  $n$  0x14 byte buffers

Our buffer will be allocated before those (we pretty much exhausted all the free chunks)



Buffer closest to our allocated buffer is the 1<sup>st</sup> one to be copied and freed

Requires a "carpet" of ~13000 contiguous 0x14 byte buffers (not too hard)




- Pros
  - 8 byte overflow only
    - The kernel pool remains pretty clean
- Cons
  - Next chunk can be:
    - A nt!\_MMFREE\_POOL\_ENTRY: we overflow a list entry
    - A busy nt!\_POOL\_HEADER
      - “Write4 on Merge with Previous” will happen on free
    - Etc
      - ⇒ Headaches
  - Higher number of packets (~300)

A horizontal banner with a red and orange background. It features a molecular structure with several hexagons and rings. The text "KNOWING YOU'RE SECURE" is written in white, uppercase letters on the left side of the banner.

KNOWING YOU'RE SECURE

# Conclusion





- Kernel still offers a nice playground for exploitation even in latest Windows versions
- Exploitation costs have increased dramatically
  - We're still working on MS08-001!



- More memory allocations are NUMA aware:
  - Up to 16 non-paged pool descriptors
    - Initial non-paged pool has separate address ranges for each node
  - Up to 16 paged pool descriptors
- This makes kernel pool overflows more complex



- How to exploit Windows kernel memory pool
  - SoBeIt, Xcon 2005
- Reliable Windows Heap Exploits
  - M. Conover, O. Horovitz, CanSecWest 2004
- Owning NonPaged pool using stealth hooking
  - mxatone, Phrack 65



KNOWING YOU'RE SECURE

Thank you!  
Questions?